

ECOBLEND OS

ecoblend.io · Future Humanity Ventures Ltd

UI/UX Design Brief

What the ECOBLEND OS interface must accomplish, and the constraints under which it may say it

Field	Value
Document ID	FHV-EB-DBR-004
Version	v1-0
Date of issue	02 August 2026
Document owner	Founder & CEO, Future Humanity Ventures Ltd
Status	Issued to design
Classification	Confidential — internal and named contractors only

Claims governance notice

This document is written under the ECOBLEND claims-governance rule: no statement may describe a capability as operational unless it is operational in code and independently verified.

Capabilities under construction are written in the future or conditional tense and carry an explicit gate reference. The terms "verified", "intelligence platform" and any AI-scoring assurance language are prohibited in external material until defects D6 and D7 are closed and builds B-02 and B-03 have passed independent verification.

1. Document control

Version	Date	Author / role	Summary of change
v1-0	02 Aug 2026	Founder & CEO, FHV	First issue. Records two live brand conflicts and proposes a token system that separates brand identity from semantic status colour.

This brief instructs the designer engaged to produce the ECOBLEND OS interface. It is not a specification of what the product does — that is FHV-EB-PRD-001 — nor of how it is built. It states what the design must accomplish, the constraints it operates under, and how the work will be judged.

2. The brief

2.1 The single job

The interface has one job: to make it effortless to see what a number rests on, and impossible to mistake absence for a low value.

Everything else — navigation, dashboards, portfolio views, exports — is instrumental to that. A design that produces a beautiful readiness dashboard and buries the calculation trace three clicks deep has failed the brief, however well it renders.

2.2 Who this is for and what they bring

Persona	What they arrive wanting	What the design must resist
Entrepreneur	A high number and a clear route to a higher one	Presenting the score as a game with a progress bar. Evidence is a claim about reality, not a scoring tactic.
Reviewer	To get through the queue	Making the model recommendation more prominent than the evidence, which would turn judgement into ratification
Investor / LP	A comparable number they can act on	Implying comparability across ventures assessed under different profiles, or implying assurance the platform does not provide
Coach	To be useful quickly	Any affordance that lets coaching activity appear to move a score
Governance reviewer	To find where the model is behaving unfairly	Aggregating away the sector-level pattern that makes unfairness visible

The asymmetry that shapes every screen

The person who produces the evidence benefits from a favourable score. The people who consume the score are harmed by an inflated one. The interface must not give the producer any control over the presentation seen by the consumer — not through layout, not through ordering, not through the selection of what appears in an export.

3. Design principles

Ref	Principle	Test of compliance
DP-1	Evidence proximity — a number is never more than one interaction from the evidence beneath it	Pick any figure on any screen. Count the interactions to its calculation trace. More than one is a defect.
DP-2	Absence is not a low value — an unevidenced dimension renders as unevidenced, never as zero, never as an empty progress bar	A screenshot of a venture with no evidence must not be capable of being read as a poor venture
DP-3	Refusal is designed, not degraded — when the system declines to produce a number it says which condition fired and what would resolve it	Every refusal condition in FTD §5.2 has a designed state with its own copy
DP-4	Honest tense — the interface never describes a capability in a tense the build does not support	No instance of "verified" or "intelligence" survives a copy audit before gate G2
DP-5	Governance is visible, not administrative — model version, profile, exclusions and overrides appear where the number appears, not in a settings panel	A score cannot be screenshotted without its governing specification appearing in the same frame

DP-5 deserves a word. The instinct is to keep versioning metadata out of the way because it clutters. The consequence is that scores circulate as screenshots stripped of the specification that gives them meaning, which is precisely the mechanism by which an assessment becomes an assurance in someone else's hands. The metadata is not clutter. It is the part that makes the number honest.

4. Brand position and two live conflicts

The existing brand guideline is a coherent identity system and is not in question. Two conflicts must nonetheless be resolved before design begins, because both will otherwise be baked into components and become expensive to unwind.

4.1 Conflict A — the palette collides with status semantics

The brand palette is built on green, amber and blue: #51AF37, #F49C13, #3A97D3, with supporting #096BA7, #2D9856, #F1C411 and #F7E6D2. Existing platform material simultaneously uses green, amber and red as venture status semantics — on track, at risk, requires intervention.

Green cannot be both the brand colour and the signal for "on track". Wherever brand green appears as chrome, a user reads reassurance that the design did not intend. In a product whose entire value proposition is that it does not flatter, this is not a cosmetic problem.

Recommended resolution

Split the token system in two. Brand tokens govern identity surfaces — masthead, marketing, export cover, empty-state illustration. Semantic tokens govern status, and are drawn from a separate, deliberately unbranded range with a distinct treatment.

Semantic status is additionally never carried by colour alone (see section 9). This is required for accessibility and, here, it is also required for meaning: the same hue is doing two jobs and one of them must be given a second channel.

4.2 Conflict B — naming is unresolved

Materials currently in circulation use "ECOBLEND OS", "ecoblend OS", "ECOBLEND IO" and "ecoblend.io" interchangeably. The business card carries "ecoblend OS" while the public platform is presented as ecoblend.io. This is decision OD-1 in the requirements document and it blocks issue of the website copy pack at v1.1.

The recommendation recorded there stands: ECOBLEND OS is the product; ecoblend.io is the domain at which it is found, and is not a second product name. Design should proceed on that basis and be revised if the founder decides otherwise, rather than waiting.

The related tagline conflict — "DESIGN | MOTION | DIGITAL" against composites and reengineering positioning — is decision OD-2 and affects the ECORACE surfaces rather than this product. It is noted here because the same designer is likely to encounter both.

5. Token system

5.1 Brand tokens

Token	Value	Use
brand/green	#51AF37	Identity surfaces only. Never status.
brand/amber	#F49C13	Identity surfaces only. Never status.
brand/blue	#3A97D3	Primary interactive colour: links, focus, selected state
brand/blue-deep	#096BA7	Headers, masthead, primary action fill
brand/green-deep	#2D9856	Identity accent at small scale
brand/yellow	#F1C411	Identity accent; not permitted on status or on any surface adjacent to a score
brand/sand	#F7E6D2	Warm neutral ground for identity and empty-state surfaces

5.2 Semantic tokens

Deliberately drawn away from the brand range so that a status signal cannot be mistaken for chrome. Every semantic token pairs with a required non-colour channel.

Token	Value	Meaning	Required second channel
state/evidenced	#1F6F52	Dimension has assessed evidence contributing to a score	Filled indicator plus explicit label
state/unevidenced	#6B7280	No evidence exists. Not a low value.	Outline indicator, hatched fill, and the words "not yet evidenced"
state/queued	#8A5A00	Submitted, awaiting independent assessment	Queue glyph plus label
state/excluded	#4B5563	Dimension excluded under a governed profile	Strike treatment plus "excluded — [profile]"
state/refused	#8B1A1A	Refusal condition fired; no score produced	Refusal panel with the specific condition named
state/stale	#7C3AED	Score predates the current model version	Version chip rendered adjacent to the figure

The stale token exists because of a specific near-term condition. Following recalculation of historic scores under the corrected mapping function, the system will hold two generations of score differing in some cases by two orders of magnitude. That state must be legible at a glance for as long as it persists, and designing it in advance is cheaper than retrofitting it during the migration.

5.3 Typography

The brand guideline specifies Avenir Next Condensed as the primary face, with Prompt as the accessible alternative and Muli, Nunito or Lato as secondary. That pairing works for identity surfaces. It does not work unmodified for a data-dense assessment interface, and the brief is explicit about why: condensed faces at small sizes reduce numeral legibility, and this product's entire currency is numerals that must not be misread.

Role	Face	Rationale
Display and identity	Avenir Next Condensed, with Prompt as web fallback	Per brand guideline. Used at scale, sparingly, on identity surfaces.
Interface body	Lato	From the approved secondary set. Neutral, wide language coverage, holds up at 14px.
Numerals and data	Lato with tabular figures enabled	Tabular figures are mandatory anywhere numbers are stacked or compared. Proportional figures in a comparison column are a legibility defect.
Code and identifiers	A monospaced face for hashes, model versions and identifiers	Content hashes and version strings must be visually verifiable character by character

5.4 Type scale

Step	Size / line	Weight	Use
Display	40 / 44	600	Identity surfaces, export covers
H1	28 / 34	600	Page title
H2	20 / 28	600	Section
Score	36 / 40	600, tabular	The score figure itself, and nothing else
Body	15 / 24	400	Prose, evidence content
Data	14 / 20	400, tabular	Table cells, component values
Meta	12 / 16	500, letterspaced 0.04em	Model version, profile, timestamps, provenance marks

The Score step is reserved. No other element in the interface may use it. This is the mechanism by which a score is always identifiable as a score and never confused with a count, a percentage or a target.

6. The signature component: the Score Block

One component carries this product's point of view. If it is right, the rest of the interface follows from it. If it is wrong, no amount of polish elsewhere recovers the brief.

6.1 Required elements

Element	Behaviour	Requirement
The figure	Score type step, tabular figures. Present only where a score legitimately exists.	FR-SCR-004
Model version chip	Adjacent to the figure, always. Never in a tooltip, never in a properties panel.	FR-SCR-002, DP-5
Profile chip	Names the governing profile and its version. Where a dimension is excluded, the chip carries the exclusion count.	FR-SCR-006
Component strip	Each dimension shown with its state token and its weight at time of calculation. Unevidenced dimensions are visible, not omitted.	DP-2, FR-SCR-004
Trace affordance	A single interaction opens the full calculation trace. Present on every instance without exception.	FR-SCR-008, DP-1
Provenance summary	Counts of self-attested against third-party attested evidence, presented as composition, not as a quality score.	FR-EVD-006
Limitations line	A persistent short statement that this is a model output under a named specification and is not an assurance, rating or recommendation.	PRD §4.2, FR-RPT-002

6.2 Prohibited behaviours

- No progress bar. A progress bar implies a defined total and a path to completion; readiness is neither.
- No traffic-light chrome around the figure. The score is a measurement, not a verdict, and the verdict belongs to a named human reviewer.
- No trend arrow without the underlying delta and the model version of both points. A movement between two model versions is not a trend.
- No aggregate "portfolio score". Averaging across profiles produces a number with no referent.
- No animation on the figure. Counting-up animation on a score is a celebration gesture, and this component is not celebrating.

6.3 The unevidenced treatment

This is the detail on which DP-2 stands or falls, and it should be the piece the designer spends most effort on. An unevidenced dimension must be visually distinct from a low-scoring one in a way that survives being screenshotted, printed in greyscale, and described aloud. The hatched outline and explicit label specified in the token table are a starting point, not a solution; the designer is asked to bring a better one.

7. Component inventory

Component	Purpose	Priority
Score Block	Section 6. The signature.	P1
Evidence card	Item with provenance mark, version history, hash, submitter, reviewer, strength and recorded basis	P1
Calculation trace panel	Components, weights, profile, mapping outputs, exclusions, arithmetic	P1
Refusal panel	Names the specific refusal condition and states what resolves it	P1
Review queue row	Assigned item with eligibility status; no self-selection affordance	P1
Decision form	Pass / hold / reject with mandatory reasoning; override path with delta capture	P1
Profile badge	Governing profile, version, exclusion count	P1
Comparison table	Cross-venture view with profile-mismatch warning surfaced before render	P2
Audit row	Actor, action, before, after, timestamp, model version	P2
Export dialog	Confirms what leaves the system and records the export	P2
Binding-constraint card	The single dimension most limiting progression, with the evidence that would close it	P2
Empty and stale states	Per section 8	P2

8. Interface copy standards

Words in this interface are load-bearing. The distinction between "not yet evidenced" and "low" is the product.

8.1 Voice

- Active voice, sentence case, plain verbs. A control names what happens when it is used, and keeps that name through the flow: the button that says "Submit evidence" produces the confirmation "Evidence submitted".
- Name things as the user recognises them, not as the system implements them. A person submits evidence; they do not create an evidence artefact record.
- Refusals explain and direct. They do not apologise and they are never vague about which condition fired.
- Empty states invite an action. They never say "no data".

8.2 Prohibited terms

Term	Status	Released at	Use instead
Verified, verification, verified evidence	Prohibited	Gate G2	Assessed, reviewed, evidence recorded
Intelligence platform, venture intelligence	Prohibited	Gate G3	Venture validation rail, assessment
AI-scored, AI-assessed, smart scoring	Prohibited	Gate G3	Scored under [model version]
Predicted, forecast, likelihood of success	Prohibited	Not planned	Readiness assessed at this date
Trusted, certified, validated venture	Prohibited	Not planned	Assessment complete; gate decision recorded
Any invitation or inducement to invest	Prohibited	On receipt of opinion under SOW-LEG-01	Descriptive presentation only

The tense rule

A capability under construction is described in the interface, in release notes and in marketing as "being built to close". Not "coming soon", which implies certainty of arrival, and not the present tense, which is simply untrue. This applies to product copy, investor decks and academic submissions alike, and the consistency is the point.

9. Accessibility

WCAG 2.1 AA is the floor and is a contractual acceptance criterion, not an aspiration. Three requirements go beyond the standard because of what this product is.

Requirement	Detail
Status never by colour alone	Every semantic token pairs with a shape, a fill treatment and a text label. This is required by the standard and independently required here, because brand green and status green would otherwise be indistinguishable.
Numerals legible at data sizes	Tabular figures throughout. Minimum 14px for any figure a decision may rest on. Condensed faces are not used for numerals at any size.
Refusals and scores announced	A refusal state and an unevidenced dimension must be conveyed to a screen reader as clearly as to a sighted user. "Not yet evidenced" is announced; it does not read as an empty cell.
Keyboard path to trace	The calculation trace is reachable by keyboard from every score instance. Visible focus throughout.
Reduced motion respected	No motion is essential to meaning anywhere in the product.

10. Responsive behaviour

Surface	Approach
Score Block	Fully responsive and never truncated. On narrow viewports the component strip stacks; the model version chip and limitations line are never the elements that drop.
Review console	Desktop-first. Assessment requires reading evidence alongside a decision form, and a narrow-viewport version that forces sequential reading changes the quality of the judgement.
Portfolio and comparison	Responsive with horizontal scroll on data tables rather than column hiding. A hidden column in a comparison is a silent omission.
Founder dashboard	Mobile-capable. Founders check status on a phone; submitting substantial evidence is a desktop activity and the design should not pretend otherwise.
Exports	A4 portrait, print-safe, greyscale-legible. Assume every export will be printed and photocopied at least once.

11. Explicitly not to be designed

Designing an interface for a capability that does not exist creates pressure to claim it exists. The following are out of scope and are listed so that their absence is a decision rather than an oversight.

- Any AI assistant, chat surface or generated-recommendation panel. No inferential capability exists.
- Any "verified" badge, trust mark, certification seal or equivalent device, in any form, at any fidelity — including in concept work.
- Predictive charts, success-likelihood indicators or forecast surfaces.
- Provenance-rail, ledger or token surfaces.
- Federated or cross-tenant benchmarking views.
- Public league tables or rankings of any kind.

Why concept work is included in this prohibition

A trust mark drawn as an exploration has a way of appearing in a deck six weeks later, at which point nobody remembers it was hypothetical. The cheapest place to enforce claims governance is before the artefact exists.

12. Deliverables and acceptance

Stage	Deliverable	Acceptance
D1	Token system: brand and semantic sets, type scale, spacing, elevation, applied to a sample surface	Brand and semantic sets demonstrably non-colliding; the conflict in section 4.1 resolved rather than restated
D2	Score Block, in all states: evidenced, unevidenced, queued, excluded, refused, stale	Passes the greyscale test and the screenshot test — an unevidenced dimension cannot be read as a low one in either
D3	The four specified flows from FHV-EB-UFD-003 at screen level, including every refusal branch	Every refusal condition in FTD §5.2 has a designed state with written copy
D4	Component library covering the P1 inventory with interaction states	Keyboard-operable; contrast verified; announced states checked with a screen reader
D5	Export template, A4, print and greyscale safe	Model version, timestamp and limitations line survive printing and photocopying
D6	Copy audit of every string produced	Zero instances of any prohibited term in section 8.2; tense rule applied throughout

12.1 The acceptance test that matters

Take a screenshot of a venture with substantial evidence in two dimensions and none in the other four. Show it to somebody who has never seen the product and ask them to describe the venture. If they say it is performing poorly, the design has failed regardless of how well it scores on every other criterion. If they say most of it has not been assessed yet, it has succeeded.

13. Open decisions

Ref	Decision	Owner	Blocking
OD-1	Canonical product name: ECOBLEND OS versus ecoblend.io	Founder	All chrome, export templates, website copy pack v1.1
OD-2	ECORACE tagline retirement	Founder	Business card and ECORACE surfaces; not this product
DB-1	Whether identity surfaces retain brand green given the semantic collision	Founder with designer	Token set D1
DB-2	Whether Avenir Next Condensed is licensed for web use at required weights, or whether Prompt becomes primary	Founder	Type scale D1
DB-3	Whether the deep teal used across the document estate is adopted into the product palette or remains document-only	Founder	Token set D1; currently the document house style and the brand guideline disagree and neither cites the other

DB-3 is a genuine inconsistency rather than a preference question. The document estate is set in deep teal; the brand guideline specifies a green, amber and blue triad and does not mention teal. Both are in active use and each was presumably right at the time. It should be resolved deliberately rather than by whichever artefact a contractor happens to open first.

End of FHV-EB-DBR-004 v1-0.